

Mixed-Mode Automatic Differentiation in



Jarrett Revels, Miles Lubin & Juan Pablo Vielma (MIT)

Introduction

Hi, I'm Jarrett

- ❑ Started writing Julia code in 2013, working on AD since 2015
- ❑ Previously worked in the Julia Group @ CSAIL under Alan Edelman
- ❑ Authored Julia's performance regressions testing facilities (*BenchmarkTools*)
- ❑ Recently started an engineering position under Juan Pablo Vielma
- ❑ Continuing to work on AD, transitioning to direct JuMP development in the fall

My Users Are Smarter Than Me

- ❑ No formal optimization background (B.S. in Physics, 2014)
- ❑ A lot of my work targets other Julia developers rather than end-users
- ❑ Downstream packages: *JuMP*, *Celeste*, *Optim*, *DifferentialEquations*, *RigidBodyDynamics*, *ValidatedNumerics*, etc...

Forward-Mode AD

Multidimensional Dual Numbers

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$

Multidimensional Dual Numbers

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f(x + \sum_{i=1}^n y_i \epsilon_i) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$

Multidimensional Dual Numbers

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f\left(x + \sum_{i=1}^n y_i \epsilon_i\right) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$



$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_i \\ \vdots \\ x_n \end{bmatrix} \rightarrow \mathbf{x}_\epsilon = \begin{bmatrix} x_1 + \epsilon_1 \\ \vdots \\ x_i + \epsilon_i \\ \vdots \\ x_n + \epsilon_n \end{bmatrix}$$

Multidimensional Dual Numbers

$$f(x + y\epsilon) = f(x) + f'(x)y\epsilon \text{ where } \epsilon \neq 0, \epsilon^2 = 0$$



$$f\left(x + \sum_{i=1}^n y_i \epsilon_i\right) = f(x) + f'(x) \sum_{i=1}^n y_i \epsilon_i \text{ where } \epsilon_i \neq 0, \epsilon_i \epsilon_j = 0$$



$$\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_i \\ \vdots \\ x_n \end{bmatrix} \rightarrow \mathbf{x}_\epsilon = \begin{bmatrix} x_1 + \epsilon_1 \\ \vdots \\ x_i + \epsilon_i \\ \vdots \\ x_n + \epsilon_n \end{bmatrix}$$



$$g(\mathbf{x}_\epsilon) = g(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g(\mathbf{x})}{\partial x_i} \epsilon_i$$

The Dual Type

```
# stack-allocated vector of partial derivatives  
using StaticArrays.SVector
```

```
# N-dimensional dual number type  
struct Dual{N,T<:Real} <: Real  
    value::T  
    partials::SVector{N,T}  
end
```

The Dual Type

```
# stack-allocated vector of partial derivatives
using StaticArrays.SVector
```

```
# N-dimensional dual number type
struct Dual{N,T<:Real} <: Real
    value::T
    partials::SVector{N,T}
end
```

```
# overload various math operations
import Base: sin, cos, -, +, *
```

```
sin(d::Dual) = Dual(sin(d.value), cos(d.value) * d.partials)
cos(d::Dual) = Dual(cos(d.value), -(sin(d.value)) * d.partials)
(-)(d::Dual) = Dual(-(d.value), -(d.partials))
(+)(a::Dual, b::Dual) = Dual(a.value + b.value, a.partials + b.partials)
(*) (a::Dual, b::Dual) = Dual(a.value * b.value,
                               b.value * a.partials + a.value * b.partials)
```

The Dual Type

```
# stack-allocated vector of partial derivatives
using StaticArrays.SVector
```

```
# N-dimensional dual number type
struct Dual{N,T<:Real} <: Real
    value::T
    partials::SVector{N,T}
end
```

```
# overload various math operations
import Base: sin, cos, -, +, *
```

```
sin(d::Dual) = Dual(sin(d.value), cos(d.value) * d.partials)
cos(d::Dual) = Dual(cos(d.value), -(sin(d.value)) * d.partials)
(-)(d::Dual) = Dual(-(d.value), -(d.partials))
(+)(a::Dual, b::Dual) = Dual(a.value + b.value, a.partials + b.partials)
(*) (a::Dual, b::Dual) = Dual(a.value * b.value,
                               b.value * a.partials + a.value * b.partials)
```

This code enables:

- sin and cos derivatives to **arbitrary order** (e.g. `Dual{M, Dual{N, T}}`)
- sin and cos derivatives over **complex number types** (e.g. `Complex{Dual{N, T}}`)
- sin and cos derivatives over **custom number types** (e.g. `Custom{Dual{N, T}}`)

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$

```
function cumprod(x)
    y = similar(x)
    if length(x) < 1
        return y
    end
    y[1] = x[1]
    for i in 2:length(y)
        y[i] = y[i-1]*x[i]
    end
    return y
end
```

Taking Jacobians using Duals

$$\mathbf{J}(\mathbf{g})(\mathbf{x}) = \begin{bmatrix} \frac{\partial g_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_1(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_1(\mathbf{x})}{\partial x_n} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \frac{\partial g_j(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_j(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_j(\mathbf{x})}{\partial x_n} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \frac{\partial g_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_m(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_m(\mathbf{x})}{\partial x_n} \end{bmatrix}$$

Taking Jacobians using Duals

$$\mathbf{J}(\mathbf{g})(\mathbf{x}) = \begin{bmatrix} \frac{\partial g_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_1(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_1(\mathbf{x})}{\partial x_n} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \frac{\partial g_j(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_j(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_j(\mathbf{x})}{\partial x_n} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \frac{\partial g_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial g_m(\mathbf{x})}{\partial x_i} & \cdots & \frac{\partial g_m(\mathbf{x})}{\partial x_n} \end{bmatrix}$$

$$\mathbf{g}(\mathbf{x}_\epsilon) = \begin{bmatrix} g_1(\mathbf{x}_\epsilon) \\ \vdots \\ g_j(\mathbf{x}_\epsilon) \\ \vdots \\ g_m(\mathbf{x}_\epsilon) \end{bmatrix} = \begin{bmatrix} g_1(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g_1(\mathbf{x})}{\partial x_i} \epsilon_i \\ \vdots \\ g_j(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g_j(\mathbf{x})}{\partial x_i} \epsilon_i \\ \vdots \\ g_m(\mathbf{x}) + \sum_{i=1}^n \frac{\partial g_m(\mathbf{x})}{\partial x_i} \epsilon_i \end{bmatrix}$$

Taking Jacobians using Duals

$$\mathbf{J}(\mathbf{g})(\mathbf{x}) = \begin{bmatrix} \boxed{\frac{\partial g_1(\mathbf{x})}{\partial x_1}} & \cdots & \boxed{\frac{\partial g_1(\mathbf{x})}{\partial x_i}} & \cdots & \boxed{\frac{\partial g_1(\mathbf{x})}{\partial x_n}} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \boxed{\frac{\partial g_j(\mathbf{x})}{\partial x_1}} & \cdots & \boxed{\frac{\partial g_j(\mathbf{x})}{\partial x_i}} & \cdots & \boxed{\frac{\partial g_j(\mathbf{x})}{\partial x_n}} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \boxed{\frac{\partial g_m(\mathbf{x})}{\partial x_1}} & \cdots & \boxed{\frac{\partial g_m(\mathbf{x})}{\partial x_i}} & \cdots & \boxed{\frac{\partial g_m(\mathbf{x})}{\partial x_n}} \end{bmatrix}$$

$$\mathbf{g}(\mathbf{x}_\epsilon) = \begin{bmatrix} g_1(\mathbf{x}_\epsilon) \\ \vdots \\ g_j(\mathbf{x}_\epsilon) \\ \vdots \\ g_m(\mathbf{x}_\epsilon) \end{bmatrix} = \begin{bmatrix} g_1(\mathbf{x}) + \sum_{i=1}^n \boxed{\frac{\partial g_1(\mathbf{x})}{\partial x_i}} \epsilon_i \\ \vdots \\ g_j(\mathbf{x}) + \sum_{i=1}^n \boxed{\frac{\partial g_j(\mathbf{x})}{\partial x_i}} \epsilon_i \\ \vdots \\ g_m(\mathbf{x}) + \sum_{i=1}^n \boxed{\frac{\partial g_m(\mathbf{x})}{\partial x_i}} \epsilon_i \end{bmatrix}$$

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)] # 3 + ε3
```

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)] # 3 + ε3
```

```
julia> cumprod(x)  
3-element Array{Dual{3,Int64},1}:  
Dual(1,1,0,0)  
Dual(2,2,1,0)  
Dual(6,6,3,2)
```

Taking Jacobians using Duals

$$\text{cumprod}\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}\right) = \begin{bmatrix} x_1 \\ x_2 x_1 \\ x_3 x_2 x_1 \\ \vdots \\ x_n x_{n-1} x_{n-2} \dots x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1 & 0 & 0 \\ x_2 x_1 & x_1 & 0 \\ x_3 x_2 & x_3 x_1 & x_2 x_1 \end{bmatrix}$$



$$\mathbf{J}(\text{cumprod})\left(\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 6 & 3 & 2 \end{bmatrix}$$

```
julia> x = [Dual(1, 1, 0, 0), # 1 + ε1  
            Dual(2, 0, 1, 0), # 2 + ε2  
            Dual(3, 0, 0, 1)] # 3 + ε3
```

```
julia> cumprod(x)  
3-element Array{Dual{3,Int64},1}:  
Dual(1, 1, 0, 0)  
Dual(2, 2, 1, 0)  
Dual(6, 6, 3, 2)
```

Perturbation Confusion

```
D = (f, x_0) -> df/dx evaluated at x_0

# nested, closed over differentiation
D(x -> x * D(y -> x + y, 1), 1)

# correct answer
d1 = D(x -> x * D(y -> x + y, 1), 1)
d1 = D(x -> x * (y -> 1)(1), 1)
d1 = D(x -> x, 1)
d1 = (x -> 1)(1)
d1 = 1
```

Perturbation Confusion

```
D = (f, x_0) -> df/dx evaluated at x_0

# nested, closed over differentiation
D(x -> x * D(y -> x + y, 1), 1)

# correct answer
d1 = D(x -> x * D(y -> x + y, 1), 1)
d1 = D(x -> x * (y -> 1)(1), 1)
d1 = D(x -> x, 1)
d1 = (x -> 1)(1)
d1 = 1
```

```
epsilon(d::Dual) = d.partials[1]

# what AD using naive dual numbers will compute
d2 = D(x -> x * D(y -> x + y, 1), 1)
d2 = D(x -> x * epsilon(x + Dual(1,1)), 1)
d2 = epsilon(Dual(1,1) *
              epsilon(Dual(1,1) + Dual(1,1)))
d2 = epsilon(Dual(1,1) * epsilon(Dual(2,2)))
d2 = epsilon(Dual(1,1) * 2)
d2 = epsilon(Dual(2,2))
d2 = 2 # != d1
```


ForwardDiff.jl

- ❑ Implements multidimensional dual numbers
- ❑ Fully stack-allocated and aggressively inlined, plays well with SIMD
- ❑ Tagging system prevents perturbation confusion and resolves nested differentiation ambiguities
- ❑ Provides an API instead of exposing dual numbers directly (e.g. `ForwardDiff.jacobian(f, x)`)
- ❑ Used by JuMP for calculating Hessian-vector products

Reverse-Mode AD



Compared to Forward-Mode AD

- ❑ Instead of propagating an input perturbation forward, we propagate an output sensitivity backwards
 - ❑ Forward-mode AD evaluates chain rule from right (inner function) to left (outer function)
 - ❑ Reverse-mode AD evaluates chain rule from left (outer function) to right (inner function)
- ❑ Main hurdle: requires a reverse-traversable computation graph
 - ❑ Graph can be defined declaratively via special objects/syntax (JuMP, TensorFlow)
 - ❑ ...or dynamically by tracking arguments/intercepting function calls (ReverseDiff, Autograd, PyTorch)
- ❑ Which mode should I use?
 - ❑ Output dimension < input dimension && input dimension >> code size? Use reverse mode
 - ❑ Output dimension > input dimension && input dimension << code size? Use forward mode
 - ❑ Output dimension = input dimension? That's tough

Code Representation

```
function f(x1, x2)  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
     $\mathbf{x}_4$  = cos( $\mathbf{x}_2$ )  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $x_1$ ,  $x_2$ )  
     $x_3$  = sin( $x_1$ )  
     $x_4$  = cos( $x_2$ )  
     $x_5$  =  $x_3$  *  $x_4$   
    # ?  
end
```


Graph Representation





Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
     $\mathbf{x}_4$  = cos( $\mathbf{x}_2$ )  
     $\mathbf{x}_5$  =  $\mathbf{x}_3$  *  $\mathbf{x}_4$   
    return  $\mathbf{x}_5$   
end
```


Graph Representation








$$x_5 = x_3 * x_4$$



$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$$x_1 = 1.0$$

$$x_2 = 1.0$$

$$x_3 = 0.0$$

$$x_4 = 0.0$$

$$x_5 = 0.0$$

$$y_1 = 0.0$$

$$y_2 = 0.0$$

$$y_3 = 0.0$$

$$y_4 = 0.0$$

$$y_5 = 0.0$$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$\rightarrow x_3 = 0.8$	$y_3 = 0.0$
$x_4 = 0.0$	$y_4 = 0.0$
$x_5 = 0.0$	$y_5 = 0.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$


$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$x_3 = 0.8$	$y_3 = 0.0$
 $x_4 = 0.5$	$y_4 = 0.0$
$x_5 = 0.0$	$y_5 = 0.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$$x_1 = 1.0$$

$$x_2 = 1.0$$

$$x_3 = 0.8$$

$$x_4 = 0.5$$

$$\rightarrow x_5 = 0.4$$

$$y_1 = 0.0$$

$$y_2 = 0.0$$

$$y_3 = 0.0$$

$$y_4 = 0.0$$

$$y_5 = 0.0$$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$


Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$x_3 = 0.8$	$y_3 = 0.0$
$x_4 = 0.5$	$y_4 = 0.0$
$x_5 = 0.4$	$y_5 = 1.0$

 $y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$		$y_1 = 0.0$
$x_2 = 1.0$		$y_2 = 0.0$
$x_3 = 0.8$	\rightarrow	$y_3 = 0.5$
$x_4 = 0.5$	\rightarrow	$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 = -0.7 \quad \leftarrow \end{aligned}$$

Numerical Results

$x_1 = 1.0$	\rightarrow	$y_1 = 0.0$
$x_2 = 1.0$		$y_2 = -0.7$
$x_3 = 0.8$		$y_3 = 0.5$
$x_4 = 0.5$		$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$y_2 += y_4 * -(\sin(x_2))$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 = 0.2 \quad \leftarrow \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 = -0.7 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	\rightarrow	$y_1 = 0.2$
$x_2 = 1.0$		$y_2 = -0.7$
$x_3 = 0.8$		$y_3 = 0.5$
$x_4 = 0.5$		$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$y_1 += y_3 * \cos(x_1)$$

$$x_4 = \cos(x_2)$$

$$y_2 += y_4 * -(\sin(x_2))$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$

ReverseDiff.jl

- ❑ Uses operator overloading to dynamically intercept and record native Julia code to an instruction tape.
- ❑ Re-recording allows for complex control flow (loops, recursion, etc.).
- ❑ Multiple dispatch + JIT + run-time type information enables compiled, specialized primitive execution methods
- ❑ Supports array primitives, linear algebraic derivative definitions, and most **AbstractArray** types
- ❑ Supports static forward/reverse passes over the tape with precomputed dispatch and preallocated instruction caches
- ❑ Leverages mixed-mode AD; intermediary scalar subgraphs can be automatically rewritten into new primitives differentiated via ForwardDiff. This includes application of scalar kernels via elementwise higher-order functions (e.g. map/broadcast).

A Few Realizations

“He must be a thorough fool who can learn nothing from his own folly.”

- A.W. Hare

Julia Is Pretty Good At This Stuff

- ❑ Multiple dispatch + JIT compilation enables seamless and precise operator overloading with essentially no performance penalties. Auto-differentiable Julia code doesn't need to "know" about ForwardDiff/ReverseDiff, as long as the code is numerically type-generic.
- ❑ Writing data-flow semantics in Julia over a Julia-represented DAG means grants efficient nested data-flow semantics for “free”.
- ❑ Since primitive definition and execution is performed via normal Julia dispatch, no magic is required to extend or create primitives.
- ❑ We can get heterogeneous device support for "free" (e.g. via GPUArrays), and hardware-specialized primitives can easily be added via dispatch

ReverseDiff For JuMP?

- ❑ Cons vs. ReverseDiffSparse:
 - ❑ ReverseDiffSparse, as the name implies, does indeed exploit Hessian sparsity
 - ❑ ReverseDiffSparse has better variable storage locality for scalar operations
- ❑ Pros vs. ReverseDiffSparse:
 - ❑ ReverseDiffSparse doesn't support array primitives
 - ❑ ReverseDiffSparse doesn't directly support native Julia code
 - ❑ ReverseDiffSparse isn't numerically type-generic
 - ❑ ReverseDiffSparse can't easily handle nested differentiation
- ❑ Takeaway: ReverseDiff is more versatile and extensible, but ReverseDiffSparse has some important performance optimizations for tackling large-scale problems

ReverseDiff For Deep Learning?

- ❑ ReverseDiff's API doesn't expose variable construction directly
 - ❑ ...though internal utilities are similar to PyTorch's/TensorFlow's exposed APIs
- ❑ ReverseDiff's dynamic recording mechanism is implemented as if its performance costs can be amortized w.r.t. whole computations
 - ❑ ...which is true for static graphs in traditional optimization
 - ❑ ...but is untrue for dynamic graphs in deep learning
- ❑ Different Graph Regimes
 - ❑ Optimization: Many nodes, computationally cheap scalar operations
 - ❑ Deep Learning: Fewer nodes, computationally expensive array operations
 - ❑ This is why ML people are cool with fully dynamic taping methods - traversal overhead is negligible

ReverseDiff For...Not AD?

- ❑ A native-Julia-to-DAG package would be generally useful outside of AD
 - ❑ Dynamic code analysis/optimization
 - ❑ Parallel operation scheduling
 - ❑ Automatic pre-allocation/memory management
 - ❑ Interval constraint programming
 - ❑ Serialization of Julia code to other DAG frameworks
- ❑ It would require generalizing ReverseDiff's taping/execution mechanisms to support arbitrary metadata propagation
- ❑ It could also enable better AD anyway (e.g. edge-pushing algorithm for sparse Hessians, which requires propagating dependency information)

Enter *Cassette.jl*

What is Cassette?

- ❑ A native-Julia-to-DAG data-flow package for forward/backward-propagating values and arbitrary metadata through pure-Julia computation graphs.
- ❑ Inspired by both deep learning and traditional optimization worlds - different representations are supported for static and dynamic graphs
- ❑ Exposes taping/execution mechanisms to downstream library authors as a hijackable processing pipeline. Cassette primitives should be easy to define and extend.
- ❑ The next version of ReverseDiff is Cassette's prototypical application

Acknowledgements

- ❑ *Juan Pablo Vielma @ MIT Operations Research Center*
- ❑ *Cosmin Petra @ Lawrence Livermore National Lab*
- ❑ *The Julia Group @ MIT CSAIL: Alan Edelman, Andreas Noack, Peter Ahrens*
- ❑ *Robin Deits @ MIT CSAIL*
- ❑ *JuMP Meetup Organizers: Miles Lubin, Chris Coey, Jennifer Challis*