

Mixed-Mode Automatic Differentiation in



Jarrett Revels, Miles Lubin & Juan Pablo Vielma (MIT)

Hi, I'm Jarrett

- ❑ Started writing Julia code in 2013, working on AD in Julia since 2015
- ❑ Authored Julia's performance regressions testing facilities (*BenchmarkTools*)
- ❑ Downstream packages: *JuMP*, *Celeste*, *Optim*, *DifferentialEquations*, *RigidBodyDynamics*, *ValidatedNumerics*, etc...
- ❑ Previously worked in the Julia Group @ CSAIL under Alan Edelman
- ❑ Recently started a new position under Juan Pablo Vielma @ MIT ORC
- ❑ Currently focused on AD, maybe transitioning to JuMP development in the fall

Last Year's Talk: ForwardDiff.jl

- ❑ Implements multidimensional dual numbers
- ❑ Fully stack-allocated and aggressively inlined, plays well with SIMD
- ❑ Provides a differentiation API instead of exposing dual numbers directly
- ❑ Tagging system prevents perturbation confusion and drives nested differentiation

This Year's Talk: More AD Stuff

- ❏ Description of reverse-mode AD
- ❏ Some reflections on ReverseDiff.jl
- ❏ A new thing I'm building that will hopefully solve most of my problems

Reverse-Mode AD



Compared to Forward-Mode AD

- ❑ Propagating input perturbation forward → propagate output sensitivity backwards
 - ❑ Forward-mode AD evaluates chain rule from right (inner function) to left (outer function)
 - ❑ Reverse-mode AD evaluates chain rule from left (outer function) to right (inner function)
- ❑ Main hurdle: requires a reverse-traversable computation graph
 - ❑ Graph can be defined declaratively via special objects/syntax (JuMP, TensorFlow)
 - ❑ ...or by running code + intercepting operations (ReverseDiff, Autograd, PyTorch)
- ❑ Which mode should I use?
 - ❑ `output_dimension > input_dimension || input_dimension << code_size` → Use forward mode
 - ❑ `output_dimension < input_dimension && input_dimension >> code_size` → Use reverse mode
 - ❑ `output_dimension ≈ input_dimension` → That's tough

Code Representation

```
function f(x1, x2)  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
     $\mathbf{x}_4$  = cos( $\mathbf{x}_2$ )  
    # ?  
end
```


Graph Representation



Code Representation

```
function f( $x_1$ ,  $x_2$ )  
     $x_3$  = sin( $x_1$ )  
     $x_4$  = cos( $x_2$ )  
     $x_5$  =  $x_3$  *  $x_4$   
    # ?  
end
```


Graph Representation





Code Representation

```
function f( $\mathbf{x}_1$ ,  $\mathbf{x}_2$ )  
     $\mathbf{x}_3$  = sin( $\mathbf{x}_1$ )  
     $\mathbf{x}_4$  = cos( $\mathbf{x}_2$ )  
     $\mathbf{x}_5$  =  $\mathbf{x}_3$  *  $\mathbf{x}_4$   
    return  $\mathbf{x}_5$   
end
```


Graph Representation








$$x_5 = x_3 * x_4$$



$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$$x_1 = 1.0$$

$$x_2 = 1.0$$

$$x_3 = 0.0$$

$$x_4 = 0.0$$

$$x_5 = 0.0$$

$$y_1 = 0.0$$

$$y_2 = 0.0$$

$$y_3 = 0.0$$

$$y_4 = 0.0$$

$$y_5 = 0.0$$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$


$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
 $x_3 = 0.8$	$y_3 = 0.0$
$x_4 = 0.0$	$y_4 = 0.0$
$x_5 = 0.0$	$y_5 = 0.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$


$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$x_3 = 0.8$	$y_3 = 0.0$
 $x_4 = 0.5$	$y_4 = 0.0$
$x_5 = 0.0$	$y_5 = 0.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$x_3 = 0.8$	$y_3 = 0.0$
$x_4 = 0.5$	$y_4 = 0.0$
$\rightarrow x_5 = 0.4$	$y_5 = 0.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$


Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	$y_1 = 0.0$
$x_2 = 1.0$	$y_2 = 0.0$
$x_3 = 0.8$	$y_3 = 0.0$
$x_4 = 0.5$	$y_4 = 0.0$
$x_5 = 0.4$	$y_5 = 1.0$

 $y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 \end{aligned}$$

Numerical Results

$x_1 = 1.0$		$y_1 = 0.0$
$x_2 = 1.0$		$y_2 = 0.0$
$x_3 = 0.8$	\rightarrow	$y_3 = 0.5$
$x_4 = 0.5$	\rightarrow	$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 = -0.7 \quad \leftarrow \end{aligned}$$

Numerical Results

$x_1 = 1.0$	\rightarrow	$y_1 = 0.0$
$x_2 = 1.0$		$y_2 = -0.7$
$x_3 = 0.8$		$y_3 = 0.5$
$x_4 = 0.5$		$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$x_4 = \cos(x_2)$$

$$y_2 += y_4 * -(\sin(x_2))$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$



Multivariable Chain Rule

$$y_i = \partial x_5 / \partial x_i$$

$$= \text{sum}(y_j * \partial x_j / \partial x_i \text{ for } j \text{ in parents}(i))$$

Derivative Outputs

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_1 &= \partial x_5 / \partial x_1 \\ &= y_1 = 0.2 \quad \leftarrow \end{aligned}$$

$$\begin{aligned} \partial f(x_1, x_2) / \partial x_2 &= \partial x_5 / \partial x_2 \\ &= y_2 = -0.7 \end{aligned}$$

Numerical Results

$x_1 = 1.0$	\rightarrow	$y_1 = 0.2$
$x_2 = 1.0$		$y_2 = -0.7$
$x_3 = 0.8$		$y_3 = 0.5$
$x_4 = 0.5$		$y_4 = 0.8$
$x_5 = 0.4$		$y_5 = 1.0$

$$x_3 = \sin(x_1)$$

$$y_1 += y_3 * \cos(x_1)$$

$$x_4 = \cos(x_2)$$

$$y_2 += y_4 * -(\sin(x_2))$$

$$x_5 = x_3 * x_4$$

$$y_3 += y_5 * x_4$$

$$y_4 += y_5 * x_3$$

ReverseDiff.jl

- ❑ Uses operator overloading to dynamically intercept and record native Julia code to an instruction tape
- ❑ Multiple dispatch + JIT + run-time type information enables compiled, specialized primitives
- ❑ Supports array primitives, linear algebraic derivative definitions, and most **AbstractArray** types
- ❑ Supports dynamic forward pass (re-recording allows for complex control flow - loops, recursion, etc.)
- ❑ Supports static forward/reverse passes over tape (precomputed dispatch + preallocated instruction caches)
- ❑ Mixed-mode AD! Scalar subgraphs automatically differentiated via ForwardDiff. Includes scalar kernels of elementwise functions (e.g. map/broadcast).

A Few Realizations

“He must be a thorough fool who can learn nothing from his own folly.”

- A.W. Hare

Julia Is Pretty Good At This Stuff

- ❑ Seamless/precise operator overloading with no performance penalty
- ❑ Target code can be mostly “AD-unaware”; just needs to be numerically type-generic.
- ❑ Primitives defined via normal Julia code - no magic for creation/extension
- ❑ Writing data-flow semantics in Julia over a Julia-represented DAG means grants efficient nested data-flow semantics for “free”.
- ❑ Heterogeneous device support for “free” (e.g. GPUArrays)

ReverseDiff For JuMP?

❑ Cons vs. ReverseDiffSparse:

- ❑ ReverseDiffSparse, as the name implies, does indeed exploit Hessian sparsity
- ❑ ReverseDiffSparse has better variable storage locality for scalar operations

❑ Pros vs. ReverseDiffSparse:

- ❑ ReverseDiffSparse doesn't support array primitives
- ❑ ReverseDiffSparse doesn't support dynamic graphs
- ❑ ReverseDiffSparse doesn't directly support native Julia code
- ❑ ReverseDiffSparse isn't numerically type-generic
- ❑ ReverseDiffSparse can't easily handle nested differentiation

- ❑ Takeaway: ReverseDiff is more versatile and extensible, but ReverseDiffSparse has some important performance optimizations for tackling large-scale problems

ReverseDiff For Deep Learning?

- ❑ ReverseDiff's API doesn't expose variable construction directly
 - ❑ ...though internal utilities are similar to PyTorch's/TensorFlow's exposed APIs
- ❑ ReverseDiff's dynamic recording mechanism writes to a static graph representation
 - ❑ ...great for recording traditional optimization graphs
 - ❑ ...not so great for recording dynamic graphs in deep learning
- ❑ Different Graph Regimes
 - ❑ Optimization: Many nodes, computationally cheap scalar operations
 - ❑ Deep Learning: Fewer nodes, computationally expensive array operations
 - ❑ This is why ML people are cool with fully dynamic taping methods - traversal overhead is negligible

ReverseDiff For...Not AD?

- ❑ A native-Julia trace-to-DAG package would be generally useful outside of AD
 - ❑ Dynamic code analysis/optimization
 - ❑ Parallel operation scheduling
 - ❑ Automatic pre-allocation/memory management
 - ❑ Interval constraint programming
 - ❑ Serialization of Julia code to other DAG frameworks
- ❑ It would require generalizing ReverseDiff's taping/execution mechanisms
- ❑ It could also enable better AD anyway (e.g. edge-pushing algorithm for sparse Hessians)

Enter *Cassette.jl*

“Multiple dispatch is dead, long live multiple dispatch!”

- Anonymous

What is Cassette?

- ❑ A native Julia execution tracer + data flow package for propagating values and arbitrary metadata through pure-Julia computation graphs.
- ❑ Inspired by both traditional optimization and deep learning worlds - different representations are supported for static and dynamic graphs
- ❑ Exposes trace interception mechanisms to downstream library authors as a hijackable processing pipeline.
- ❑ The next version of ReverseDiff is Cassette's prototypical application
- ❑ ***Doesn't rely on argument type propagation to intercept function calls!***

Multiple Dispatch Is Dead...

```
# we'll define primitives for this on the next slide
```

```
struct Interceptor{T,N} <: AbstractArray{T,N}
```

```
    data::AbstractArray{T,N}
```

```
end
```


Multiple Dispatch Is Dead...

```
# we'll define primitives for this on the next slide
```

```
struct Interceptor{T,N} <: AbstractArray{T,N}  
    data::AbstractArray{T,N}  
end
```

```
# primitives defined on `Interceptor` will just call this
```

```
struct Intercepted{F} <: Function  
    func::F  
end
```

```
unwrap(x) = x
```

```
unwrap(i::Interceptor) = i.data
```

```
unwrap(i::Intercepted) = i.func
```

```
(i::Intercepted{F})(args...) = (println("called $F");  unwrap(i)(unwrap.(args)...))
```

Multiple Dispatch Is Dead...

```
const AMBIGUOUS_TYPES = [subtypes (AbstractArray) ...]
```

```
#### 1-arg primitive --> 1 method #####
```

```
Base.f(x::Interceptor) = Intercepted(f)(x)
```

Multiple Dispatch Is Dead...

```
const AMBIGUOUS_TYPES = [subtypes(AbstractArray)...]
```

```
#### 1-arg primitive --> 1 method #####
```

```
Base.f(x::Interceptor) = Intercepted(f)(x)
```

```
#### 2-arg primitive --> ~50 methods! #####
```

```
Base.f(x::Interceptor, y::Interceptor) = Intercepted(f)(x, y)
```

```
for T in AMBIGUOUS_TYPES
```

```
    Base.f(x::Interceptor, y::T) = Intercepted(f)(x, y)
```

```
    Base.f(x::T, y::Interceptor) = Intercepted(f)(x, y)
```

```
end
```

Multiple Dispatch Is Dead...

```
const AMBIGUOUS_TYPES = [subtypes(AbstractArray)...]
```

```
#### 1-arg primitive --> 1 method #####
```

```
Base.f(x::Interceptor) = Intercepted(f)(x)
```

```
#### 2-arg primitive --> ~50 methods #####
```

```
Base.f(x::Interceptor, y::Interceptor) = Intercepted(f)(x, y)
```

```
for T in AMBIGUOUS_TYPES
```

```
    Base.f(x::Interceptor, y::T) = Intercepted(f)(x, y)
```

```
    Base.f(x::T, y::Interceptor) = Intercepted(f)(x, y)
```

```
end
```

```
#### 3-arg primitive --> ~2000 methods!!! #####
```

```
Base.f(x::Interceptor, y::Interceptor, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
for T in AMBIGUOUS_TYPES
```

```
    Base.f(x::Interceptor, y::Interceptor, z::T) = Intercepted(f)(x, y, z)
```

```
    Base.f(x::Interceptor, y::T, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
    Base.f(x::T, y::Interceptor, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
    for S in AMBIGUOUS_TYPES
```

```
        Base.f(x::Interceptor, y::T, z::S) = Intercepted(f)(x, y, z)
```

```
        Base.f(x::T, y::Interceptor, z::S) = Intercepted(f)(x, y, z)
```

```
        Base.f(x::T, y::S, z::Interceptor) = Intercepted(f)(x, y)
```

```
    end
```

```
end
```

Multiple Dispatch Is Dead...

```
const AMBIGUOUS_TYPES = [subtypes(AbstractArray)...]
```

```
#### 1-arg primitive --> 1 method #####
```

```
Base.f(x::Interceptor) = Intercepted(f)(x)
```

```
#### 2-arg primitive --> ~50 methods #####
```

```
Base.f(x::Interceptor, y::Interceptor) = Intercepted(f)(x, y)
```

```
for T in AMBIGUOUS_TYPES
```

```
    Base.f(x::Interceptor, y::T) = Intercepted(f)(x, y)
```

```
    Base.f(x::T, y::Interceptor) = Intercepted(f)(x, y)
```

```
end
```

```
#### 3-arg primitive --> 200 methods #####
```

```
Base.f(x::Interceptor, y::Interceptor, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
for T in AMBIGUOUS_TYPES
```

```
    Base.f(x::Interceptor, y::Interceptor, z::T) = Intercepted(f)(x, y, z)
```

```
    Base.f(x::Interceptor, y::T, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
    Base.f(x::T, y::Interceptor, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
    for S in AMBIGUOUS_TYPES
```

```
        Base.f(x::Interceptor, y::T, z::S) = Intercepted(f)(x, y, z)
```

```
        Base.f(x::T, y::Interceptor, z::S) = Intercepted(f)(x, y, z)
```

```
        Base.f(x::T, y::S, z::Interceptor) = Intercepted(f)(x, y, z)
```

```
    end
```

```
end
```

GROSS

...Long Live Multiple Dispatch!

```
#### AST pruning pseudocode #####
```

```
function code_info_with_intercepted_calls(Tuple{F,A,B,C...})  
    # 1. Get `CodeInfo` for signature `Tuple{F,A,B,C}`  
    # 2. Walk through SSA-form AST and wrap all calls with `Intercepted`  
    # 3. Return the modified `CodeInfo`  
end
```

```
#### `Trace` function wrapper #####
```

```
struct Trace{F,world} <: Function  
    func::F  
    Trace(func::F) where {F} = new{F,get_world_counter()}(func)  
end
```

```
@generated function (t::Trace{F,world})(args...) where {F,world}  
    return code_info_with_intercepted_calls(F, args...)  
end
```

...Long Live Multiple Dispatch!

```
function f(x)
    a = one(eltype(x))
    b = 100 * a
    result = zero(eltype(x))
    for i in 1:length(x)
        result += b * (a - x[i])
    end
    return result
end
```

...Long Live Multiple Dispatch!

```
function f(x)
  a = one(etype(x))
  b = 100 * a
  result = zero(etype(x))
  for i in 1:length(x)
    result += b * (a - x[i])
  end
  return result
end
```

↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓

```
function (::Trace{typeof(f)})(x)
  a = Intercepted(one)(Intercepted(etype)(x))
  b = Intercepted(*) (100, a)
  result = Intercepted(zero)(Intercepted(etype)(x))
  for i in Intercepted(UnitRange)(1, Intercepted(length)(x))
    result = Intercepted(+)(result, Intercepted(*) (b, Intercepted(-)(a,
      Intercepted(getindex)(x, i))))
  end
  return result
end
```


...Long Live Multiple Dispatch!

```
function f(x)
  a = one{eltype}(x)
  b = 100 * a
  result = zero{eltype}(x)
  for i in 1:length(x)
    result += b * (a - x[i])
  end
  return result
end
```

↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓↓

```
function (::Trace{typeof(f)})(x)
  a = Intercepted{one}(Intercepted{eltype}(x))
  b = Intercepted{*}(100, a)
  result = Intercepted{zero}(Intercepted{eltype}(x))
  for i in Intercepted{UnitRange}(1, Intercepted{length}(x))
    result = Intercepted{+}(result, Intercepted{*}(b, Intercepted{-}(a,
      Intercepted{getindex}(x, i))))
  end
  return result
end
```

```
(i::Intercepted)(ws::MyWrapper...) = MyWrapper.(unwrap(i)(unwrap.(ws)...))
```

...Long Live Multiple Dispatch!

- ❑ No need for target functions to be type generic
- ❑ No need to define an ungodly number of methods per primitive
- ❑ No need to define new number/array/etc. types just to propagate metadata or hijack execution flow
- ❑ Hijack behavior can be overloaded via normal Julia dispatch of downstream function wrappers
- ❑ In the future, we can also wrap SSA-form control flow instructions

The Future

- ❑ Finish + document + test + release Cassette (targeting Julia v0.7)
- ❑ Replace ForwardDiff/ReverseDiff with new Cassette-based packages
- ❑ Replace ReverseDiffSparse → new backend for JuMP
 - ❑ Locality + sparse Hessian optimizations for Cassette graphs
- ❑ Evangelize Cassette for other regimes

Acknowledgements

- ❑ *Juan Pablo Vielma, Miles Lubin @ MIT Operations Research Center*
- ❑ *Cosmin Petra @ Lawrence Livermore National Lab*
- ❑ *The Julia Group @ MIT CSAIL: Alan Edelman, Andreas Noack, Peter Ahrens*
- ❑ *Jameson Nash, Mike Innes @ Julia Computing (...and everybody else there as well!)*
- ❑ *Simon Danisch @ GPUArrays, Inc*
- ❑ *Robin Deits @ MIT CSAIL*